

PocketData Benchmark

Naveen, Sankar, Saravanan, Sathish

What did we do last week?

- Started extracting data from phone log.
- 93b46e40acbf1167ab0ad421b73761f16b74b562 1426486358667 1426486358667.0 2015-03-16
06:12:38.667574 13274 13274 I SQLite-Query-PhoneLab
{"Counter":0,"LogFormat":"1.0","AppName":"Google Play
Books","Action":"APP_NAME","PackageName":"com.google.android.apps.books"}
- 93b46e40acbf1167ab0ad421b73761f16b74b562 1426477052477 1426477052477.1 2015-03-16
03:37:32.477583 13274 13274 I SQLite-Query-PhoneLab
{"Counter":561,"LogFormat":"1.0","Time":243698,"Arguments":"null","Results":"SELECT value FROM ScheduledTaskProto
ORDER BY sortingValue ASC, insertionOrder ASC","Action":"SELECT","Rows returned":3}
- {"Counter":27,"LogFormat":"1.0","Time":129375,"Results":"SQLiteProgram: INSERT INTO
carriers(bearer,authtype,carrier_enabled,protocol,mmsproxy,roaming_protocol,numeric,mcc,type,mmsc,password,mvno_m
atch_data,mvno_type,name,server,mnc,apn,user,mmsport) VALUES
(?,?)","Action":"INSERT","Arguments(hashCoded)":"0,-1,1231,2343,-
1630285132,2343,47743056,49679,-697368471,155249537,117478,0,0,1755004508,42,1537,1954370557,117478,1784,"}
- {"Counter":5,"LogFormat":"1.0","Time":149843,"Arguments":"null","Results":"PRAGMA
table_info(name)","Action":"SELECT","Rows returned":0}

How did we parse it?

- Switched to Java from Python
 - Java is faster. Need to parse more than a GB
 - Better support with jsqlparser
- Extracted features like number of projected columns, table count etc.,
 - Made use of jsqlparser visitor pattern
- Modified jsqlparser to allow PRAGMA and INSERT or REPLACE queries (for future).

Challenges faced

- It took 3.5 hours to parse and extract log files.
- Future : To run parallel for users.


Method Name Filter (Contains)		
Hot Spots - Method	Self Time [%]	Self Time
net.sf.jsqparser.parser.CCJsSqlParser. Statement ()		10,643,050 ms (88.2%)
net.sf.jsqparser.parser.CCJsSqlParser.<init> (java.io.Reader)		639,304 ms (5.3%)
org.json.simple.parser.JSONParser. parse (String)		191,458 ms (1.6%)
java.io.Writer. append (CharSequence)		128,450 ms (1.1%)
java.io.BufferedReader. readLine ()		119,369 ms (1%)
java.lang.reflect.Field. get (Object)		86,316 ms (0.7%)
edu.ub.tbd.service.PersistanceFileService. write (edu.ub.tbd.entity.AbstractEntity)		71,163 ms (0.6%)
java.lang.String. split (String)		33,168 ms (0.3%)
java.sql.Timestamp. toString ()		31,389 ms (0.3%)
edu.ub.tbd.parser.LogParser. parseSingleLogFile (String)		27,965 ms (0.2%)

Challenges [contd...]

- Unable to find app names for 32 million queries.
 - Something wrong with our logic
- Coming up with features without data available.

Results

Rows Returned Vs Time Taken (Micro secs)



Next week plan

- Play around with the extracted data to come up with interesting features.
- Extract more features from the log files.